

Discourse Coherence, Reference Grounding and Goal Oriented Dialogue

Baber Khalid[†] Malihe Alikhani^{†*} Michael Fellner* Brian McMahan[†] Matthew Stone^{†*}

[†]Rutgers Computer Science Department

*Rutgers Center for Cognitive Science

{firstname.lastname}@rutgers.edu

Abstract

Prior approaches to realizing mixed-initiative human–computer referential communication have adopted information-state or collaborative problem-solving approaches. In this paper, we argue for a new approach, inspired by coherence-based models of discourse such as SDRT (Asher and Lascarides, 2003), in which utterances attach to an evolving discourse structure and the associated knowledge graph of speaker commitments serves as an interface to real-world reasoning and conversational strategy. As first steps towards implementing the approach, we describe a simple dialogue system in a referential communication domain that accumulates constraints across discourse, interprets them using a learned probabilistic model, and plans clarification using reinforcement learning.

1 Introduction

Referential communication tasks have long served benchmark problems for human–computer dialogue systems (Heeman and Hirst, 1995), and they remain a topic of current investigation (Monroe et al., 2017). One reason is the flexibility and robustness of human joint activity in referential communication (Clark and Wilkes-Gibbs, 1986). Another is the complexity of bootstrapping and reconciling wide-coverage models of semantics in dialogue, on the one hand, and effective models of collaboration under uncertainty, on the other.

Computational work on dialogue has generally integrated collaboration with semantic processing using act-based models (Traum, 1994; Larson and Traum, 2000; Rich et al., 2001; DeVault, 2008; Galescu et al., 2018). On these approaches, problem-solving actions are recognized indirectly from utterance content using intention recognition, and thereby integrated into a hierarchical model of the dialogue’s open goals and plans. While these

approaches are powerful and successful, it remains challenging to scale up the necessary models of dialogue structure and intention recognition. For example, system builders have to map out an inventory of transitions that interlocutors can use to update the problem-solving state and have to describe how those transitions are signaled by utterances with specific kinds of form and meaning.

This paper advocates an alternative approach to link the theory of collaborative activity to the practice of dialogue system design, inspired by coherence-based approaches to dialogue semantics such as SDRT (Asher and Lascarides, 2003; Lascarides and Asher, 2009). On a coherence-based approach, collaboration plays out at the level of the discourse as a whole, rather than at the level of the individual utterance. In particular, semantic interpretation attaches utterance content directly into an evolving discourse structure, from which the commitments of interlocutors can be derived as a knowledge graph; it is coherence-based interpretation, rather than plan recognition, that builds a hierarchical structure for the dialogue. Collaborative reasoning comes in later, to connect the coherent dialogue content to its implications for interlocutors’ goals, plans and problem solving. In particular, semantic evaluation operations, trained by reinforcement learning or supervised learning from generic resources, ground the knowledge graph in the task context and underwrite system strategy.

We begin this paper with a high-level perspective on this approach, highlighting the relevance of discourse-level interpretation for collaborative reasoning and sketching the possible benefits of eliminating step-by-step models of the task contributions of utterances in favor of holistic task-based evaluations of discourse content. As a first step in implementing this approach, we describe a proof-of-concept system that distinguishes arbitrary color patches from alternatives, as in Monroe

et al. (2017). We emphasize our extensive (but not complete) use of self-supervised and reinforcement learning methods in constructing the system, and report preliminary results showing the effectiveness of a baseline version of the system in interaction with crowd workers. This experience suggests that a coherence-based approach offers a promising alternative direction for systematic efforts to scale up computational approaches to flexible, mixed-initiative dialogue.

2 Motivation

The key linguistic insight of our work is that referential communication depends on semantic grounding operations that apply at the level of the discourse as a whole, operations that draw not only on the explicit commitments of interlocutors but also on those signalled implicitly. We begin by showing how these insights are captured by coherence approaches to discourse representation, particularly SDRT (Asher and Lascarides, 2003; Lascarides and Asher, 2009), and motivating their importance for referential communication tasks. We argue further that coherence-based approaches create novel, and potentially attractive engineering trade-offs in capturing these phenomena in interactive systems.

As the empirical starting point for our work, we use the Colors in Context (CIC) dataset of Monroe et al. (2017). They asked participants to talk about items in a visual display using a free-form chat interface. On each round of interaction, one human subject, designated the director, was provided privately with a target item from a randomized display and tasked with instructing the other human subject, designated the matcher, to click on the correct item. Examples from this dataset are shown in Figure 1–4. In line with the longstanding findings of empirical work on referential communication since Clark and Wilkes-Gibbs (1986), both director and matcher assume responsibility in these dialogues to make sure that the team members arrive at a mutually satisfactory description of the target patch and complete the task successfully.

2.1 Coherence

Coherence theory (Hobbs, 1979, 1985; Kehler, 2001; Asher and Lascarides, 2003) characterizes conversation through a small number of qualitative relations that link each utterance to an evolving representation of the dialogue. Linking an utterance to a related segment can create a sibling (coordina-




Figure 1: An example from the Colors in Context (CIC) dataset (Monroe et al., 2017) of the director and matcher coordinating so that the matcher can click on the correct color patch (x_0).

tion) or a child (subordination), generating a hierarchical structure. In some approaches to discourse coherence, including SDRT, contributions can have multiple attachments, so discourse structure is a complex graph, not just a tree.¹ This seems to be particularly important for situated conversation (Hunter et al., 2018), including talk about visual contexts like the displays of the CIC dataset.

Coherence relations come in a number of qualitatively different varieties. Some relations signal inferences at the level of domain entities: *Narrative* is a coordinating relation involving descriptions of successive events; *Expansion* is a subordinating relation involving the extended description of a single entity. Some relations are best described in terms of propositions: *Parallel* is a coordinating relation that presents two propositions as analogous; *Explanation* is a subordinating relation connecting a fact to the generalization that explains it. A final class of relations involve meta-level relationships that involve the goals or circumstances of utterances. *Correction* in dialogue is a coordinating relationship of this kind; *Clarification Request* is a subordinating one. Coherence theories postulate substantial divergences between these three different classes of relations. However, relations of any kind suffice to make the conversation coherent and to underwrite inferences about the content of the conversation as a whole, including the commitments of the different speakers.

Schlangen and Lascarides (2003) argue for the suitability of using coherence to characterize the interpretation of the fragmentary utterances—such as those in Figure 1—that are ubiquitous in referential communication, and Stone and Lascarides (2010) show how representations of coherence and speaker commitment can underwrite probabilistic inferences about interlocutors’ shared understand-

¹Hobbs (1985) and Kehler (2001) seem to focus on tree-like discourse structures, for example.

ing in dialogue. We begin by reviewing, extending and deepening those formal accounts. Alternative approaches to these issues are of course possible (Ginzburg, 2012, among others), and we are certainly not in a position to argue against other approaches at this stage of our work. However, we do argue that coherence theory offers an elegant, integrated framework that addresses the structural, semantic and inferential challenges of referential communication in a way that extends to other kinds of dialogue interaction as well.

2.2 Coherence and Dialogue Structure

Coherence theory offers a simple and natural way to represent the STRUCTURE of interactions such as that in Figure 1. Let’s look at this interaction in detail. It’s useful to segment the utterances into basic moves that establish reference, commit to propositions and contribute open questions. This lets us describe the dialogue as an action u_1 that refers to the target with the description *the very dull red*; a question u_2 whether the target is *lighter* and a question u_3 whether the target is *darker*, combined via an alternation relation into segment σ_4 contributing a complex question; and a contribution u_5 that the target is *lighter*. These fit together into a hierarchical structure:

- u_1 discourse initial
- σ_4 *Alternative*(u_2, u_3)
subordinate to u_1
Clarification(u_1, σ_4)
- u_5 coordinate to σ_4
Answer(σ_4, u_5)
and *Answer*(u_2, u_5)

The hierarchical structure here follows from the definitions of the relations involved—*Clarification*, *Alternative*, *Answer*—and is built incrementally as moves attach to the discourse. Nevertheless, it supports the key inferences you’d expect for a hierarchical discourse structure, including the possibility for moves to attach high and return to an ongoing topic after a subsidiary inquiry, and the incoherence that interlocutors typically feel when topics are disjointed or interleaved. Moreover, it delivers these inferences as part of general model of dialogue that can be relevant for other tasks such as planning successful dialogue. For example, recognizing terms of the matcher’s question *lighter or darker* are parallel alternatives (which might be specified or learned as a typical feature of effective clarification questions).

2.3 Coherence and Semantic Reference

Coherence theory also offers a sophisticated approach to the SEMANTICS of referential communication dialogues. Like all referential communication, the dialogue of Figure 1 is about its contextual situation: the director’s utterances work to identify patch x_0 . It’s common to think of such situated reference in terms of the compositional semantics and pragmatics of individual expressions, as in Kripke’s (1978) philosophical account of speaker reference or the words-as-classifiers paradigm of Kennington and Schlangen (2015). However, researchers as early as Luperfoy (1992) have argued that the management of situated reference in the presence of anaphoric relations and partial information requires representing and reasoning about the content of the discourse as a whole.

Figure 2 shows how the problem comes up in the CIC dataset. In cases such as that shown in Figure 2, directors frame difficult reference problems not just by describing the target but also by giving contrasting descriptions of relevant alternatives. Subsequent utterance can identify objects in context by exploiting these established expressions.

Coherence theory might describe this dialogue in a fine-grained way via the composition of four moves by the director (for now we ignore the contribution of *i see*). The director begins with three descriptions that introduce discourse referents into an evolving discourse model: Such inferences are the bread-and-butter of the dynamic approaches to discourse semantics on which coherence approaches such as SDRT are based; see Dekker (2012). We will represent this as three fragments u_1 , u_2 and u_3 , related by a *Parallel* relation into a segment σ_4 . This segment is subordinate to the follow up *the pinkish one* and attaches to it via a *Background* relation; *the pinkish one* also attaches to u_3 by a *Presupposition Resolution* relation through which the definite noun phrase is resolved anaphorically to the discourse referent introduced by u_3 . The result is an integrated representation of the speaker’s




Figure 2: Semantic grounding needs to apply at the level of the discourse to handle anaphoric relations.

commitments in the conversation that features just three discourse referents, each of which is constrained to pick out a different entity in the visual scene, and one of which is designated as giving the target referent for the task.

Note how coreference constraints, disjoint reference constraints, and visual grounding work together at the level of the discourse in interpreting the director's contribution. Evidently, the description *pinkish one* on its own leaves substantial uncertainty about the intended referent; indeed, x_2 might be a more likely match than x_0 . However, the implicit constraint that the target can be distinguished from another patch that's *purple* and another that's *reddish* significantly reduces the possibility of misinterpretation. Any constraint-satisfaction approach to ambiguity resolution will handle this, whether logical (Mellish, 1985) or probabilistic (Gorniak and Roy, 2005). But what's to prevent the original uncertainties to resurface when we interpret a new token of the description *the pinkish one*? The natural way is to follow the coherence-based account, and invoke a discourse-level inference that interprets a second occurrence as dependent on (and therefore referentially identical to) the first, in virtue of an anaphoric link.

2.4 Coherence and Common Ground

Coherence itself is vital for calculating an integrated representation of dialogue content (Lascarides and Asher, 2009). Content is often contributed and acknowledged implicitly, in virtue of the coherence relations that attach utterances to the discourse. This is easily regimented in coherence theory. Lascarides and Asher (2009) give a formal model that exploits coherence to identify the commitments that interlocutors have made over the course of a coherent discourse. Such calculations can be crucial to tracking mutual understanding in dialogue (Stone and Lascarides, 2010), as we can see even in the CIC domain.




Figure 3: Clarification questions implicate the matcher's implicit acceptance of the possibility of grounding the director's description in the visual scene.

Take clarification requests, as in Figures 1 and 3. Even though clarification requests flag particular aspects of interpretation as incomplete, they implicitly accept the descriptions they target as appropriate to the visual scene. Thus, the matcher of Figure 1 seems willing to accept x_0 or x_1 as a very dull red; that of Figure 3 seems willing to accept x_0 or x_1 as a sky blue. In particular, the dialogue of Figure 3 illustrates how these implicit commitments can inform interlocutors' subsequent coordination: the utterance *more greenish* must be understood to select from the sky blue candidates that have been mutually agreed; x_3 may well be the *more greenish* patch in the display as a whole.

Figure 4 shows what happens when no such implicit commitments can be inferred. Since the matcher's response completely rejects the director's initial description, the director proceeds with an alternative referential strategy—exclusion of the two other patches—that stands independently of the content contributed by the initial utterance.

2.5 Implementing Coherence

In this paper, we explore coherence not just as a theoretical model for dialogue but as an architecture for implementing dialogue systems. Our overarching vision is that utterances in dialogue should be interpreted with a grammatically-specified meaning that's resolved against the linguistic context and attached to an evolving discourse structure. This meaning need not specify abstract updates to the state of the interaction, such as grounding acts or problem-solving acts; it just specifies what's required to track the content of the discourse. It is discourse-level inferences that link speakers' commitments to the real-world task context and underwrite the prediction of the outcomes of conversational strategies based on system experience.

We can contrast this architecture with problem-




Figure 4: Implicit commitments contrast with outright rejections. Here, the director's followup strategy of excluding distractors seems uninformed by expectations or conclusions the matcher might have drawn from the initial utterance.

solving approaches that capture semantic grounding and contributions to collaborative activity by reasoning about individual utterances. Compared to a discourse-level approach, an utterance-level architecture is likely to require a finer-grained integration of system capabilities, correspondingly finer-grained data resources—and, given current methods—substantially more handcrafted representation and data annotation.

3 System Architecture

Realizing the vision we have sketched in Section 2 represents a long-term project for dialogue systems design. In particular, it requires us to adapt linguistic processing models so that utterance interpretations contribute to a structured logical form for conversation; it requires us to adapt reasoning and planning frameworks so they apply at the level of discourse content. Will it be possible to make the changes involved?

We think it will. In the rest of this paper, we argue for the practical feasibility of a coherence-based approach to dialogue architecture, using our initial experience implementing a dialogue system using the design principles articulated in Section 2 that successfully plays the referential game investigated by Monroe et al. (2017). As described in Section 3.1, we used the development dialogues of the CIC dataset to design linguistic processing models that interpret utterances as coherent updates to discourse-level representations of logical form. Then, as described in Section 3.2, we combined machine-learned models of lexical semantics and referential interpretation with rule-based models of compositional semantics to characterize the referential ambiguities of these discourse representations. Finally, as described in Section 3.3, we characterize system abilities in terms of a set of high-level strategies for contributing to conversation, and implement a reinforcement-learning approach to selecting those strategies based on their discourse-level effects. These three interacting reasoning models work together to operationalize our approach to referential communication as an empirically-based, discourse-level collaborative process.

3.1 Handling Utterances

Our approach to utterance interpretation is based on a handcrafted context free grammar that compositionally interprets utterances in terms of corresponding formal contributions to dialogue.

The basic units of content in the system are color predications drawn from the XKCD lexicon presented in McMahan and Stone (2015). Using the inference methods described in Section 3.2, we can interpret these predications descriptively, simply as saying what a color is like, or referentially, as identifying the unique patch that fits its content. Compositional operations allow these predications to be combined by logical operations into complex descriptions. Contributions to logical form can introduce discourse referents and characterize them in terms of descriptions; these contributions are defined recursively in terms of coherence relations, which are attached into the overarching discourse structure. Given a discourse structure, we can traverse it to determine the commitments of the director, which involve a set of discourse referents, a sequence of descriptive characterizations of those discourse referents, and inequalities among discourse referents. One of these discourse referents is distinguished as representing the target of the referential communication task.

Our logical form language offers the following moves at the level of discourse: *Identify*, *Distinguish*, *Ask-Clarification*, *Confirmation*, *Rejection*. *Identify*(x, D) commits the speaker that description D identifies discourse referent x ; for example, *the green one* corresponds to a move *Identify*($T, green$) (for the distinguished target T). *Distinguish*(x, y) commits the speaker that discourse referents x and y are different; for example *not the dark blue* corresponds to two related moves *Distinguish*(T, A); *Identify*($A, dark\ blue$). In this simple implementation, *Ask-Clarification* proposes a description of the target object, to which nobody is committed. *Confirmation* (*Rejection*) commit the speaker that this description fits (does not fit) the target.

Every utterance contributes a node which attaches to existing node(s) in the discourse graph. Initially, the discourse graph consists of a single node introducing the target discourse referent. With each new description of target by the director we add a new node in the graph which attaches to most recent node by the director. Whenever a new clarification question comes, it adds a new node in the graph which attaches to the initial node. A confirmation or rejection attaches to the most recent node representing a clarification question and the most recent node by the director.

We handle the idiosyncratic syntax and tele-

graphic structure of color words, acknowledgments and other frequent contributions in CIC utterances using a domain-specific syntactic specification, developed by analyzing the development set of the CIC corpus.² We illustrate our approach with a small excerpt from our context-free grammar (CFG) in Table 1.

Category	Expansions	
$\langle S \rangle$	CP	$NegP$
$\langle CP \rangle$	$ADJ : CLR$	CLR
$\langle ADJ \rangle$	<i>grassy</i>	<i>super</i>
$\langle NegP \rangle$	$NEG : CP$	$NEG : ADJ$
$\langle NEG \rangle$	not	
$\langle CLR \rangle$	Lux Lexicon	

Table 1: A small portion of the CFG developed for our dialogue system’s interpretation engine (‘:’ denotes string concatenation).

We extend our basic grammar to a probabilistic CFG (PCFG) by automatically inducing the relevant weights from the CIC train split. In particular, we obtain all parse trees of the 15285 utterances in the CIC train split, determine the commitments the parse trees contribute to the dialogue, and then weight the trees by the probability that those commitments converge on the actual target of the interaction using the models of Section 3.2. The resulting weighted expansions are used to estimate the probabilities associated with our CFG rules.

We use Earley’s algorithm for CFG parsing and A* search to obtain the most likely parse tree given a PCFG. To handle incomplete parses, we used the *first-set* of grammar variables to find the next position where parser could resume to find a sub-parse tree.

Table 2 shows the coverage of our parser over the train split of CIC dataset. The performance shows that our lexicon and grammar rules were able to capture the vast majority of words and structures used by human interlocutors in this setting.

To evaluate parser accuracy, we looked at the first utterances in each dialogue of the CIC test set, and focused on just those cases where the parser returned at least one complete parse. We took the most likely parse tree and computed its semantic interpretation and evaluation. If the actual target of

²For example, when analyzing the data, it became apparent that the fact that whether a given occurrence of color name figured syntactically as either a noun or an adjective had no bearing on the meaning derived for the utterance as a whole.

Parse Status	Rate
Complete Parse	.8000
One NOPP	.1329
Two NOPPs	.0353
Three or More NOPPs	.0020
No Parse	.0283

Table 2: Parsing coverage on CIC training data, including complete parses, non-overlapping partial parses spanning segments of the input (NOPPs), and complete parse failure.

that interaction had the largest posterior probability of the candidate patches, we counted this as a successful item. Overall, the success rate was 88.9%. Even though this evaluation method does not confirm the correctness of the parse itself, we believe that it provides strong evidence for the broad accuracy of our parser.

3.2 Semantic Evaluation

Given these representations of discourse content, we define a semantic evaluation operation to link that discourse content to the visual world. The input is the structure of speaker commitments. The output is a posterior distribution over candidate color patches for the target. (Although in principle our logical form language allows us to compute posteriors over color patches for any of the discourse referents evoked in the dialogue, we focus here on the typical case where we only need to track the distinguished discourse referent for the target.) In other words, semantic evaluation defines a probability distribution

$$P(x_i | w^t, w^{t-1}, \dots, w^0, C)$$

where w^t represents the descriptive constraint on the target associated with the t th speaker commitment, x_i ranges over the candidate color patches in the scene, and C indicates the dependency of the calculation on a specific visual context, in this case given by three color patches.

For atomic descriptions from the XKCD lexicon, we use the models developed in McMahan and Stone (2020), which are available at <https://go.rutgers.edu/ugycmlb0>. These models are learned in two steps directly from naturally-occurring data. The first step is to train a model of color term meaning from Randall Munroe’s crowd-sourced collection of free text descriptions of color patches, as curated by McMahan and Stone (2015).

The next step is to use cognitive modeling and a latent variable model to estimate the prevalence of different speaker strategies for using color meanings to construct identifying descriptions from the CIC dataset. The result is a probability distribution

$$P(w_k|x_i, C)$$

describing the likelihood that a human speaker will use the term w_k to identify a target x_i in the context C . This is readily converted to $P(x_i|w_k, C)$ by Bayes's theorem (on the assumption that each of the targets is equally likely *a priori*).

To extend atomic probabilities to complex predication, we use logical definitions. For example, for a conjoined predication $w^{t1} \wedge w^{t2}$ we have:

$$P(x_i|w^{t1} \wedge w^{t2}) = P(x_i|w^{t1})P(x_i|w^{t2})$$

For disjunction, we get $P(x_i|w^{t1} \vee w^{t2}) =$

$$1 - ((1 - P(x_i|w^{t1}))(1 - P(x_i|w^{t2})))$$

by de Morgan's rule.

Finally, we treat the commitments of the speaker as providing independent evidence about the identity of the target. Again under the assumption that all targets are equally likely *a priori*, this gives

$$P(x_i|w^t, w^{t-1}, \dots, w^0, C) \propto \prod_{i=0}^t P(x_i|w^i, C)$$

3.3 Learning Data-Driven Decisions

Act-based approaches to collaboration plan contributions to conversation based on formal models of the preconditions and effects of moves in conversation: the system can reason backward from conversational goals to the actions it might use to achieve them. In our approach, the formal effects of conversational actions are defined exclusively by the commitments of interlocutors, as represented in discourse structure, which offers a comparatively thin notion of the dynamics and goals of conversation. We therefore need a different approach to reason about what actions to take next. Our approach is reinforcement learning, which is based on characterizing the outcomes that count as successes in dialogue and anticipating the course of conversational interaction.

A reinforcement learner represents its decision making as a Markov Decision Process (MDP), defined by states S , actions A , a transition function T , rewards R , and discount factor γ . It iteratively

improves its estimates of the expected payoffs for an action a in a state s based on experience, and ultimately selects the action whose estimated payoff is the highest. In particular, using data from the Colors in Context dataset, we train a reinforcement learner to approximate the state-action value function under a decision-making policy π , otherwise known as the Q-function, which is given as:

$$Q^\pi(s, a) = r(a) + \gamma * E[\max_{a' \in A} Q^\pi(s', a')]$$

We use a feature-based representation of the dialogue state and function approximation to drive a uniform, general estimate of the Q-function.

Developing this approach involves establishing formal models for the actions, reward, state and effects of conversational actions. Our system chooses from two possible actions in each turn, *Ask_Clarification* and *Select*. The description associated with the *Ask_Clarification* move is determined by the referential goal communication model of McMahan and Stone (2020) applied to the most likely candidate referent; we generate the most likely expression not previously used in the dialogue. The *Select* action, meanwhile, picks the most likely color patch μ under the posterior defined in Section 3.2.

We represent the state s^t using features so as to record the history of the dialogue and the current information available about the target:

$$s_t = \{\forall x_i P(x_i|w^t, \dots, w^0, C) : [a^0, \dots, a^t]\}$$

The prior actions in the dialogue are represented as a vector using a one-hot encoding.

The reward function r is designed to create a tradeoff between dialogue length and task success; good dialogues are both short and effective. In particular, if the correct target is t

$$\begin{aligned} r(\text{Select}) &= 0.3 \text{ if } \mu = t \\ r(\text{Select}) &= -0.5 \text{ if } \mu \neq t \\ r(\text{Ask_Clarification}) &= -0.1 \end{aligned}$$

After 15 turns the dialogue ends in failure with a reward of -1 and no further actions are possible.

In order to anticipate the long-term effects of system choices, we model the responses that the system elicits. This is known as a user simulation. Initially, the simulation identifies the target using a sampled description controlled by a parameter P_x which represents the probability the user will succeed in generating an identifying description. With probability P_x , the user presents a description

that’s merely true of the target patch using the semantic model of McMahan and Stone (2020), but with probability $1 - P_x$ the user’s description is sampled from their model of identifying descriptions. If the system clarifies, the user simulation confirms if the most likely referent of the distribution is the target, it responds with a *Confirmation*; otherwise it responds with a *Rejection* and a new sampled description.

4 Analysis and Results

In Section 3, we presented our initial approach to characterizing semantic evaluation and collaborative reasoning at the level of discourse structure. We’ve seen that we can characterize utterances effectively in this framework. Now we show further that the semantic evaluation operations and collaborative reasoning models are effective.

4.1 Human Evaluation

We ran a human evaluation experiment with a simple rule-based baseline system that uses the semantic evaluation models of Section 3.2 to decide on dialogue strategy. This baseline system went for a *Select* action whenever the probability of most likely referent was greater than 0.95 and chose to ask a clarification request otherwise.

Protocol. This study was conducted with the approval of our human subjects review committee. We recruited 50 subjects through Amazon Mechanical Turk. Participants were all US citizens, gave written consent, and were compensated at an estimated rate of USD 15 an hour. Each subject played 4 games with the system out of which two were selected from the CLOSE condition of Monroe et al. (2017), one from their FAR condition and one from their SPLIT condition. (According to our sensitivity power analysis, with a sample size of 176 conversations we would be able to detect effect sizes as small as 0.820 with a power and significance level of 90%.) After each trial, subjects were asked to rate the performance of the system on a scale of 0 to 5 and leave us feedback.

Results. The overall rate of successful trials is 72.6% and the average overall rating of the system is 3.7. In understanding the effectiveness of discourse-level semantic operations, we are particularly interested in the course of dialogues where information accumulates over time. There are actually 92 cases where the system simply chooses to




			
	x_0	x_1	x_2
director:	Pink		
matcher:	The hot pink?		
director:	Hot pink		
matcher:	[SELECT]		

Figure 5: An example conversation of a human director interacting with our system using a baseline matcher policy.

select on the first turn, picking the correct target in 76 of them (over 82%). In the 123 trials where the system asks for clarification after the first turn, just 78 led to successful identification—a rather disappointing success rate of just over 63%. However, the problem seems to be the users’ interpretation of system utterances, not the system’s integration of content across the discourse. System clarification questions include descriptions that the utterance-level model of McMahan and Stone (2020) predicts will identify a particular color patch. If that patch is in fact the target of this round, we’d expect users to confirm; if that patch is not the target of this round, we’d expect users to reject. However, users’ answers align with the answers we would expect them to give based on the true target only 48% of the time. Truly effective systems will need to ask clarification questions that lead to more reliable answers.

Nevertheless, Figure 5 shows an example conversation of the user with our system that shows how the use of clarification does sometimes lead to a natural and effective conversation in this task. We are releasing the data collected for this evaluation.³

4.2 RL Experiments

We ran computer experiments in our RL framework to demonstrate that our approach converges to sensible policies for contributing to a collaborative interaction.

Method. We estimated action outcomes using the deep Q learning (DQN) algorithm of Mnih et al. (2015). The Q function is approximated using linear function θ_p called POLICY NET. Another linear function θ_t also known as TARGET NET is used to predict the future value of a state s . A delta δ is calculated as below and the gradient loss of this

³<https://github.com/baber-sos/Discourse-Coherence-Reference-Grounding-and-Goal-Oriented-Dialogue>

error is back-propagated to learn the Q function.

$$\delta = Q_{\theta_p}^{\pi}(s, a) - (r(a) + \gamma * \max_{a' \in A} Q_{\theta_t}^{\pi}(s', a'))$$

Training is done by sampling a mini-batch of size 24 from an experience replay memory of size 200 with a circular buffer, calculating the δ using the Huber loss and then back-propagating using the Adam optimizer. Also, instead of updating θ_t at each iteration using back-propagation, it gets updated by copying the weights of θ_p every 20 iterations. We set the discount factor $\gamma = 1$, learning rate $\alpha = 0.001$ and weight decay parameter $\lambda = 0.01$ during the training process. We train our RL framework for 4000 dialogue simulations on color samples from CIC train split. The model is then tested on 400 simulations over unseen color samples.




Figure 6: Progression of training loss over 4000 simulated conversations with color samples for CIC dataset train split.

Results. We first report experiments with $P_x = 0.7$, which show that the DQN algorithm is effective at optimizing the expected results of system decisions during training. Figure 6 shows the trend of the training loss and Figure 7 presents the reward progression during training. Both show the system consistently improving over time. Table 3 shows the success rate over test data after each epoch. Average loss over testing data after training is 0.038. Both measures indicate that the model is generalizing well to unseen examples. Ultimately, the simulated success rate is 92% for unseen color patches at test time.




Figure 7: Trend line of average training reward.

	Successes	Failures
1st Epoch	378	22
2nd Epoch	369	31

Table 3: Success vs Failures over Test Set when following the learned policy π . The minor difference in the number of successes is because of imprecision in simulation.

Our next experiments contrast the policies learned with $P_x = 0.7$ against policies learned with $P_x = 0.4$. The differences show that the RL setup is able to learn strategies that respond appropriately to differences in simulated user behavior. Figure 8 shows histograms of model choices at the first turn, as a function of the probability value associated with the most likely referent for the target. Recall that the options are *Selection*, which commits the system to a specific referent, and *Clarification*, which aims to get more information. We can see that the model trained with $P_x = 0.7$ —with users that are less successful in identifying the target in the first turn—chooses to ask clarification questions more often than the model trained with $P_x = 0.4$. Thus as it was expected, the model learns to select the color patch only when it is almost certain that it can match the director’s description with the target.

We observe that the model trained with $P_x = 0.7$ learned to ask 66.7% more clarification questions at probability values ≥ 0.8 . Similarly, it tends to select a patch 5.5% less at the probability values ≥ 0.8 . This also showcases how P_x controls the trade-off between clarification and selection.




Figure 8: A comparison of the number of *Clarification* actions and *Select* actions chosen by the models.

5 Related Work

The role of the dialogue manager in dialogue systems is to interact in a natural way to help the users complete the task that the system is designed to support (Walker, 2000). Reinforcement learning has been traditionally successful for learning optimized yet flexible strategies in dialogue (Frampton and Lemon, 2006; Manuvinakurike et al., 2017; Walker, 2000; Zhang et al., 2018). To automatically learn dialogue management strategies from dialogue data, these models rely on simulated users, as we do; see Schatzmann et al. (2006) for a review. To learn interaction policies, we follow the approach of Fang et al. (2013); She and Chai (2017),

who use a simulated human model built from the dataset (they learn grounded verb semantics interactively in situated human–robot dialogue).

Several projects have studied deep RL for continuous state (and action) spaces, generally in robotics applications (Lillicrap et al., 2016, among others). Nevertheless, to the best of our knowledge, our work is the first that studies the effectiveness of applying deep RL models for learning clarification strategies in dialogue. In fact, however, because these methods make decisions based on the values of continuous parameters, they offer similar flexibility to approaches to dialogue management based on partially-observable MDPs, which are widely used in dialogue management (Williams and Young, 2007; Young et al., 2013). A comparison of partially observable MDPs to deep RL in this domain is left for future work.

Our findings provide new insights into ways that interlocutors employ disambiguation strategies while suggesting the potential discourse approaches in collaborative problem solving in dialogue. The closest work to ours is Schlangen (2004), who presents a model of possible causes for requesting clarifications in dialogue while relating models of communication (Clark, 1996; Allwood, 2000) to the discourse semantics of Asher and Lascarides (2003). Our work goes beyond this approach by providing tools for scaling it up, including by learning semantic models and dialogue strategies.

6 Conclusion

In this paper, we have described the rationale and design principles for using coherence in mixed-initiative referential communication dialogues, and described a preliminary implementation. We showed that we could engineer significant coverage of natural human–human dialogue utterances and strategies, drawing systematically on a range of machine learning methods, and how we could fine-tune the system’s purposeful behavior using reinforcement learning drawing on cognitively inspired user simulations. At the same time, we suggested that the system’s representations and architecture promise better coverage of anaphorically-mediated real-world reference and dialogues that mix agreement, partial agreement and explicit disagreement.

Some gaps issues remain to offer a complete and convincing demonstration of our framework. We need better performance evaluations, along with

closer links between the reinforcement learning models and the system's interaction with people, ultimately naive users with an intrinsic interest in communicating with the system. We need to substantiate the capacity of the architecture to support mixed initiative by learning strategies where the system acts as the director, not just as the matcher. We need to explore more systematic processing of discourse in context, including a robust and general treatment of anaphora and attachment. And we need to start explicitly drawing on general methods for parsing and DRS construction. Though we would have preferred to use domain-independent tools and resources, it remains a general challenge to build high-precision linguistic resources for natural dialogue. These diverse open problems offer exciting opportunities for future research.

Acknowledgement

The research presented here is supported by NSF Awards IIS-1526723 and CCF-19349243. We thank the Semdial reviewers for helpful comments, and the Mechanical Turk annotators for their contributions.

References

- Jens Allwood. 2000. An activity based approach to pragmatics. In H. Bunt and W. Black, editors, *Abduction, Belief and Context in Dialogue: Studies in computational pragmatics*, pages 47–80. John Benjamins Publishing Company.
- Nicholas Asher and Alex Lascarides. 2003. *Logics of Conversation*. Cambridge University Press, Cambridge, UK.
- Herbert H Clark. 1996. *Using language*. Cambridge university press.
- Herbert H. Clark and Deanna Wilkes-Gibbs. 1986. Referring as a collaborative process. *Cognition*, 22(1):1–39.
- Paul J. E. Dekker. 2012. *Dynamic Semantics*. Number 91 in Studies in Linguistics and Philosophy. Springer, Berlin.
- David DeVault. 2008. *Contribution Tracking: Participating in Task-Oriented Dialogue under Uncertainty*. Ph.D. thesis, Rutgers University, New Brunswick, NJ.
- Rui Fang, Changsong Liu, Lanbo She, and Joyce Chai. 2013. Towards situated dialogue: Revisiting referring expression generation. In *Proceedings of the 2013 Conference on Empirical Methods in Natural Language Processing*, pages 392–402.
- Matthew Frampton and Oliver Lemon. 2006. Learning more effective dialogue strategies using limited dialogue move features. In *Proceedings of the 21st International Conference on Computational Linguistics and the 44th annual meeting of the Association for Computational Linguistics*, pages 185–192. Association for Computational Linguistics.
- Lucian Galescu, Choh Man Teng, James Allen, and Ian Perera. 2018. *Cogent: A generic dialogue system shell based on a collaborative problem solving model*. In *Proceedings of the 19th Annual SIGdial Meeting on Discourse and Dialogue*, pages 400–409, Melbourne, Australia. Association for Computational Linguistics.
- Jonathan Ginzburg. 2012. *The interactive stance*. Oxford University Press, Oxford.
- Peter Gorniak and Deb Roy. 2005. *Probabilistic grounding of situated speech using plan recognition and reference resolution*. In *Proceedings of the 7th International Conference on Multimodal Interfaces, ICMI 05*, page 138143, New York, NY, USA. Association for Computing Machinery.
- Peter A. Heeman and Graeme Hirst. 1995. *Collaborating on referring expressions*. *Computational Linguistics*, 21(3):351–382.
- Jerry R. Hobbs. 1979. Coherence and coreference. *Cognitive Science*, 3(1):67–90.
- Jerry R. Hobbs. 1985. On the coherence and structure of discourse. Technical report, Center for the Study of Language and Information, Stanford University.
- Julia Hunter, Nicholas Asher, and Alex Lascarides. 2018. *A formal semantics for situated conversation*. *Semantics and Pragmatics*.
- Andrew Kehler. 2001. *Coherence, Reference and the Theory of Grammar*. CSLI, Stanford.
- Casey Kennington and David Schlangen. 2015. *Simple learning and compositional application of perceptually grounded word meanings for incremental reference resolution*. In *Proceedings of the 53rd Annual Meeting of the Association for Computational Linguistics and the 7th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 292–301, Beijing, China. Association for Computational Linguistics.
- Saul Kripke. 1978. Speaker's reference and semantic reference. In Peter A. French, Theodore E. Uehling, Jr., and Howard K. Wettstein, editors, *Contemporary Perspectives in the Philosophy of Language*, pages 6–27. University of Minnesota Press, Minneapolis.
- Staffan Larsson and David Traum. 2000. Information state and dialogue management in the TRINDI dialogue move engine toolkit. *Natural Language Engineering*, 6:323–340.

- Alex Lascarides and Nicholas Asher. 2009. Agreement, disputes and commitment in dialogue. *Journal of Semantics*, 26(2):109–158.
- Timothy P. Lillicrap, Jonathan J. Hunt, Alexander Pritzel, Nicolas Heess, Tom Erez, Yuval Tassa, David Silver, and Daan Wierstra. 2016. [Continuous control with deep reinforcement learning](#). In *4th International Conference on Learning Representations, ICLR 2016, San Juan, Puerto Rico, May 2-4, 2016, Conference Track Proceedings*.
- Susann Luperfoy. 1992. [The representation of multimodal user interface dialogues using discourse pegs](#). In *30th Annual Meeting of the Association for Computational Linguistics*, pages 22–31, Newark, Delaware, USA. Association for Computational Linguistics.
- Ramesh Manuvinakurike, David DeVault, and Kalliroi Georgila. 2017. Using reinforcement learning to model incrementality in a fast-paced dialogue game. In *Proceedings of the 18th Annual SIGdial Meeting on Discourse and Dialogue*, pages 331–341.
- Brian McMahan and Matthew Stone. 2015. [A Bayesian model of grounded color semantics](#). *Transactions of the Association for Computational Linguistics*, 3:103–115.
- Brian McMahan and Matthew Stone. 2020. Analyzing speaker strategy in referential communication. In *The 21st Annual Meeting of the Special Interest Group on Discourse and Dialogue (SIGDIAL 2020)*.
- Christopher S. Mellish. 1985. *Computer Interpretation of Natural Language Descriptions*. Ellis Horwood, Chichester, UK.
- Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A Rusu, Joel Veness, Marc G Bellemare, Alex Graves, Martin Riedmiller, Andreas K Fidjeland, Georg Ostrovski, et al. 2015. Human-level control through deep reinforcement learning. *Nature*, 518(7540):529–533.
- Will Monroe, Robert X.D. Hawkins, Noah D. Goodman, and Christopher Potts. 2017. [Colors in context: A pragmatic neural model for grounded language understanding](#). *Transactions of the Association for Computational Linguistics*, 5:325–338.
- C. Rich, C. L. Sidner, and N. Lesh. 2001. COLLAGEN: applying collaborative discourse theory to human-computer interaction. *AI Magazine*, 22:15–25.
- Jost Schatzmann, Karl Weilhammer, Matt Stuttle, and Steve Young. 2006. A survey of statistical user simulation techniques for reinforcement-learning of dialogue management strategies. *Knowledge Engineering Review*, 21(2):97–126.
- David Schlangen. 2004. Causes and strategies for requesting clarification in dialogue. In *Proceedings of the 5th SIGdial Workshop on Discourse and Dialogue at HLT-NAACL 2004*, pages 136–143.
- David Schlangen and Alex Lascarides. 2003. [The interpretation of non-sentential utterances in dialogue](#). In *Proceedings of the Fourth SIGdial Workshop of Discourse and Dialogue*, pages 62–71.
- Lanbo She and Joyce Chai. 2017. Interactive learning of grounded verb semantics towards human-robot communication. In *Proceedings of the 55th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 1634–1644.
- Matthew Stone and Alex Lascarides. 2010. Coherence and rationality in dialogue. In *Proceedings of the 14th SEMDIAL Workshop on the Semantics and Pragmatics of Dialogue*, pages 51–58, Poznan.
- David R Traum. 1994. *A Computational Theory of Grounding in Natural Language Conversation*. Ph.D. thesis, University of Rochester.
- Marilyn A Walker. 2000. An application of reinforcement learning to dialogue strategy selection in a spoken dialogue system for email. *Journal of Artificial Intelligence Research*, 12:387–416.
- Jason D Williams and Steve Young. 2007. Partially observable markov decision processes for spoken dialog systems. *Computer Speech & Language*, 21(2):393–422.
- Steve Young, Milica Gai, Blaise Thomson, and Jason D. Williams. 2013. Pomdp-based statistical spoken dialog systems: A review. *Proceedings of the IEEE*, 101(5):1160–1179.
- Jiaping Zhang, Tiancheng Zhao, and Zhou Yu. 2018. Multimodal hierarchical reinforcement learning policy for task-oriented visual dialog. In *Proceedings of the 19th Annual SIGdial Meeting on Discourse and Dialogue*, pages 140–150.